James Wiesebron

jwiesebron@pm.me | South Lake Tahoe, CA | jameswiesebron.com | github.com/jwbron

Engineer focused on AI agents and productivity tooling for knowledge work: adoption and training, agent guardrails, evals and cost-efficient AI usage. Backed by ten years of platform and data engineering, including work with sensitive student, health and financial data. Writes about runtime alignment and agent safety at jameswiesebron.com.

Experience

Khan Academy: Software Engineer **2022–2026**

- Founding engineer on Khan Academy's AI Developer Experience team. Wrote its charter, set best practices for AI tools, and taught engineers and non-engineers to use them.
- Owned the design and launch of an AI code reviewer on GitHub Actions and the Claude Agent SDK that became default-on for every pull request in Khan's frontend and backend monorepos. About 12 specialist sub-agents with findings validated in code, threads resolved as authors addressed them, and releases gated by an eval bench and about 2,000 tests. It replaced a managed review bot, covering every PR and push at about one fifth of that bot's estimated cost for the same coverage.
- Built coding-agent harnesses with OS-level sandboxing, credential brokering and scoped permissions. Enforced required workflow checks in code and routed agent-reported tooling failures to a shared backlog and Jira, with human triage for ambiguous reports.
- Built eval infrastructure with repeatable A/B runs, regression monitoring, transcript capture and cost tracking by model. Owned AI spend, evaluated vendors and cut costs by making usage more efficient rather than reducing it.
- Built and operated autoscaling, self-hosted GitHub Actions runners on Cloud Run that halved compute costs for CI. Drove Terraform adoption with a reusable framework and migrated services from App Engine, Compute Engine and Kubernetes Engine to Cloud Run.
- Wrote RFCs and ADRs, gave internal talks, interviewed and onboarded engineers, and led incidents on call. Led the migration of Khanmigo support feedback to Reforge analytics with a TPM.

Asana: Data Engineer (L5) **2021–2022**

- Designed and built a metrics analysis framework and dashboards that helped data scientists and analysts investigate changes, forecast consistently across teams, and maintain company-wide visibility into metric definitions.
- Developed feature-store infrastructure and tooling for data scientists to iterate on ML models, including experiment tracking, data quality assessment, and scalable model evaluation.

Rally Health: Senior Software Engineer II **2019–2021**

- Led a team through the rewrite of a healthcare search engine on Scala and Elasticsearch during its migration from Elasticsearch 2 to 6, and rebuilt a high-volume batch ETL pipeline as distributed streaming with Kafka.

Intuit: Senior Software Engineer **2018–2019**

- Designed and built a backend web service using event-driven, distributed message processing and scoring for the QuickBooks rewards program.

Capital One: Software Engineer, Technology Development Program **2016–2018**

- Built a Scala and Akka orchestration framework for containerized ML models behind credit card application and fraud decisions, partnering with data scientists (full-time, 2017–2018).

Skills

AI adoption and training, vendor evaluation, agent harnesses and guardrails, evals, Python, Go, TypeScript, GCP, Cloud Run, GitHub Actions, Terraform, Kubernetes, Kafka, Claude Agent SDK, CI/CD.

Education

University of North Carolina at Chapel Hill: Bachelor's degree in Computer Science, May 2017. Teaching assistant for Data Structures.