James Wiesebron

jwiesebron@pm.me | South Lake Tahoe, CA | github.com/jwbron

Backend and platform engineer focused on developer productivity and reliable agent workflows. Experience owning platforms from design through operations, leading infrastructure migrations, and working with internal users to keep automated work useful and aligned with human intent.

Experience

Khan Academy: Software Engineer **2022–2026**

- Owned the design, rollout, and operation of an AI code-review harness built on github actions and the claude agent SDK. Orchestrated specialized reviewers, validated findings, and scoped follow-up reviews to subsequent changes. Replaced a commercial bot at about one fifth of the cost while improving review quality and depth and incorporating engineers' feedback.
- Built and operated a custom cloud run-based system for autoscaling, self-hosted github actions runners that halved compute costs. Owned rollout, incident response, monitoring and recovery.
- Migrated services and jobs from app engine, compute engine and kubernetes engine to cloud run jobs, functions and services. Managed google cloud (GCP) infrastructure with terraform, including a reusable framework, engineering-wide architecture decisions and service-team migrations.
- Built coding-agent harnesses with OS-level sandboxing, credential brokering, and scoped permissions. Enforced required workflow checks in code and connected agent-reported tooling failures to a shared backlog and jira, with human triage for ambiguous reports.
- Built evaluation infrastructure with repeatable A/B runs, regression monitoring, transcript capture, and cost tracking by model. Used it to guide harness and model choices, distinguishing model failures from problems in test setup and scoring.
- Led the technical migration of support requests and feedback to reforge analytics for teams supporting khanmigo, in partnership with a TPM. Advised on the internal replacement.

Asana: Data Engineer (L5) **2021–2022**

- Designed and built a metrics analysis framework and dashboards that helped data scientists and analysts investigate changes, forecast consistently across teams, and maintain company-wide visibility into metric definitions.
- Developed feature-store infrastructure and tooling for data scientists to iterate on ML models, including experiment tracking, data quality assessment, and scalable model evaluation.

Rally Health: Senior Software Engineer II **2019–2021**

- Designed and built a high-volume streaming data pipeline introducing distributed processing with kafka, deployed using docker.
- Led a team through the rewrite and migration of a healthcare search engine and web service built on Scala and elasticsearch, including monolith decomposition and performance optimization.

Intuit: Senior Software Engineer **2018–2019**

- Designed and built a backend web service using event-driven, distributed message processing and scoring for the quickbooks rewards program.

Capital One: Software Engineer, Technology Development Program **2016–2018**

- Built a dependency-graph-based orchestration framework for containerized ML models, partnering with data scientists to translate research requirements into infrastructure (full-time, 2017–2018).
- Developed a full-stack s3 bucket-management tool for an internal open-source dataset metadata search engine during a 2016 internship.

Skills

Python, Go, TypeScript, GCP, cloud run, github actions, terraform, kubernetes, kafka, claude agent SDK, pi, developer tooling, CI/CD, agent orchestration, evals, sandboxed execution.

Education

University of North Carolina at Chapel Hill: Bachelor's degree in Computer Science, May 2017.